



# Overlay Networks

Reading: 9.4

COS 461: Computer Networks

Spring 2006 (MW 1:30-2:50 in Friend 109)

Jennifer Rexford

Teaching Assistant: Mike Wawrzoniak


<http://www.cs.princeton.edu/courses/archive/spring06/cos461/>

# Goals of Today's Lecture



- Routing overlays
  - Experimental versions of IP (e.g., 6Bone)
  - Multicast (e.g., MBone and end-system multicast)
  - Robust routing (e.g., Resilient Overlay Networks)
- Types of peer-to-peer networks
  - Directory-based (e.g., original Napster design)
  - Unstructured (e.g., Gnutella, Kazaa, BitTorrent)
  - Structured (e.g., distributed hash tables)
- Challenges in peer-to-peer
  - Legal issues, free riding, fast response to queries, peers coming and going over time, reliability, security, ...

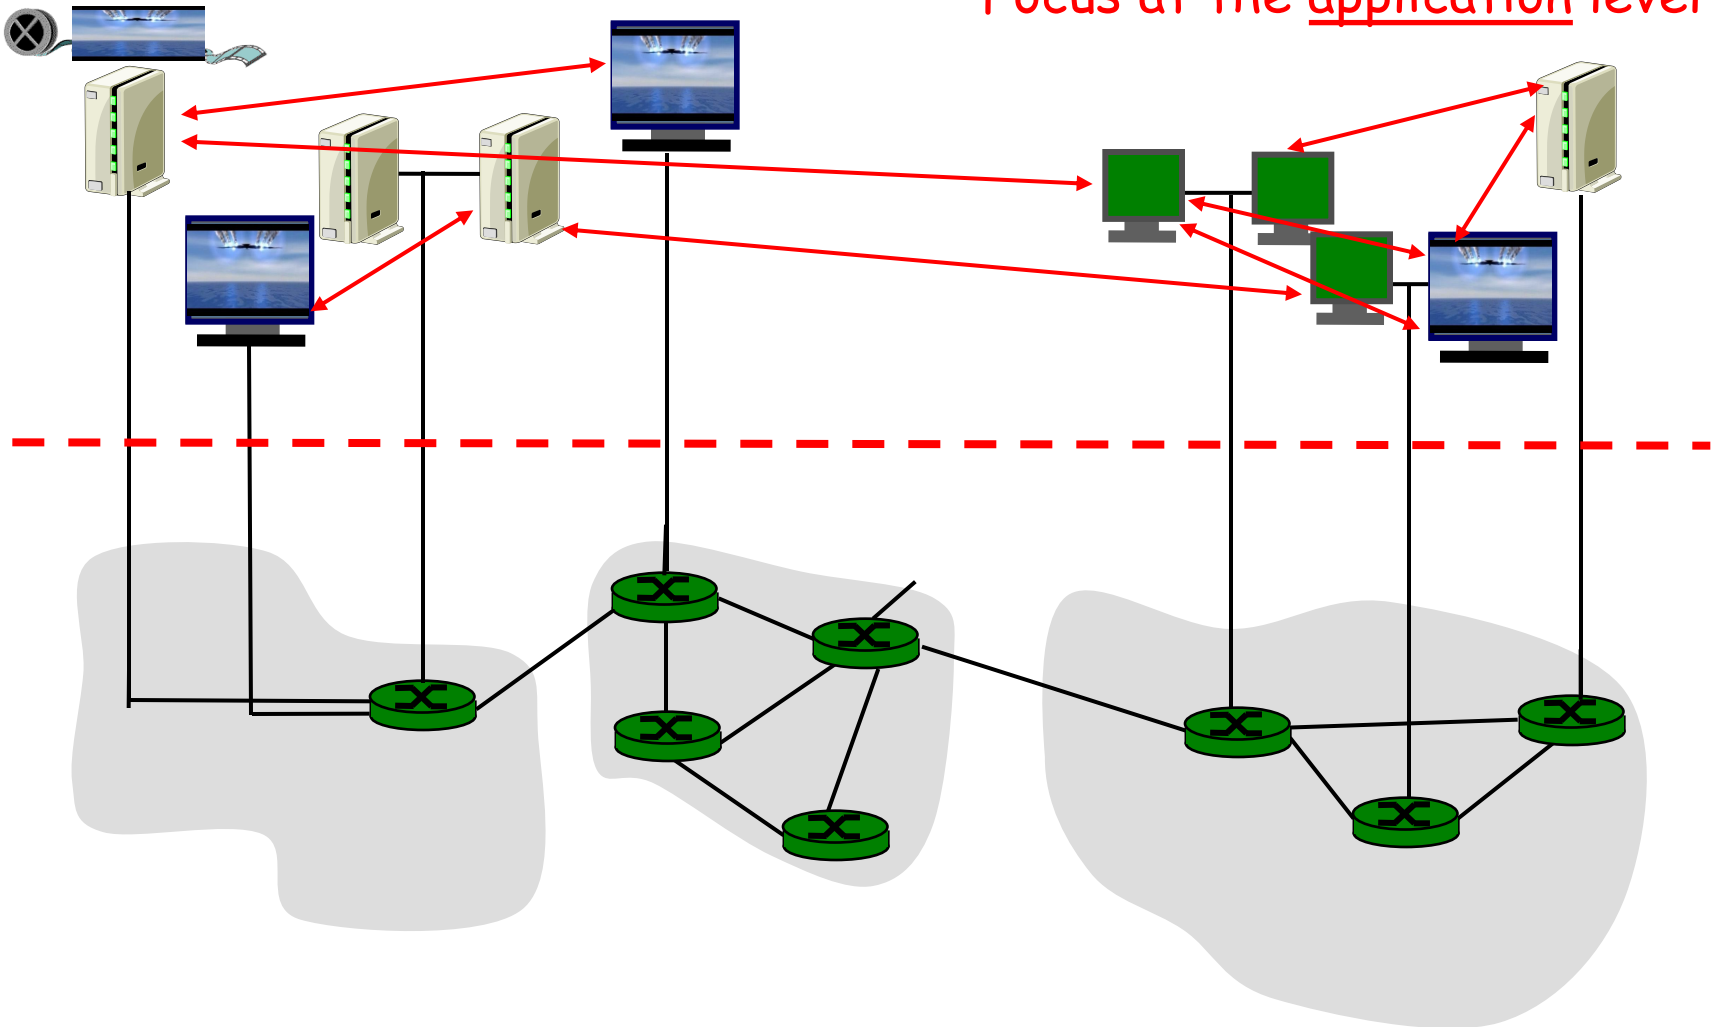
# Overlay Networks



# Overlay Networks



Focus at the application level



# Overlay Networks



- A logical network built on top of a physical network
  - Overlay links are tunnels through the underlying network
- Many logical networks may coexist at once
  - Over the same underlying network
  - And providing its own particular service
- Nodes are often end hosts
  - Acting as intermediate nodes that forward traffic
  - Providing a service, such as access to files
- Who controls the nodes providing service?
  - The party providing the service (e.g., Akamai)
  - Distributed collection of end users (e.g., peer-to-peer)


# Routing Overlays

- Alternative routing strategies
  - No application-level processing at the overlay nodes
  - Packet-delivery service with new routing strategies
- Incremental enhancements to IP
  - IPv6
  - Multicast
  - Mobility
  - Security
- Revisiting where a function belongs
  - End-system multicast: multicast distribution by end hosts
- Customized path selection
  - Resilient Overlay Networks: robust packet delivery

# IP Tunneling



- IP tunnel is a virtual point-to-point link
  - Illusion of a direct link between two separated nodes



- Encapsulation of the packet inside an IP datagram
  - Node B sends a packet to node E
  - ... containing another packet as the payload


# 6Bone: Deploying IPv6 over IP4



Logical view:



Physical view:






# Secure Communication Over Insecure Links




- Encrypt packets at entry and decrypt at exit
- Eavesdropper cannot snoop the data
- ... or determine the real source and destination



# Communicating With Mobile Users




- A mobile user changes locations frequently
  - So, the IP address of the machine changes often
- The user wants applications to continue running
  - So, the change in IP address needs to be hidden
- Solution: fixed gateway forwards packets
  - Gateway has a fixed IP address
  - ... and keeps track of the mobile's address changes



# MBone: IP Multicast

- Multicast

- Delivering the same data to many receivers
- Avoiding sending the same data many times




- IP multicast

- Special addressing, forwarding, and routing schemes
- Not widely deployed, so MBone tunneled between nodes

# End-System Multicast


- IP multicast still is not widely deployed
  - Technical and business challenges
  - Should multicast be a *network-layer* service?
- Multicast tree of end hosts
  - Allow end hosts to form their own multicast tree
  - Hosts receiving the data help forward to others



# RON: Resilient Overlay Networks



**Premise:** by building application overlay network,  
can increase performance and reliability of routing



# RON Can Outperform IP Routing



- IP routing does not adapt to congestion
  - But RON can reroute when the direct path is congested
- IP routing is sometimes slow to converge
  - But RON can quickly direct traffic through intermediary
- IP routing depends on AS routing policies
  - But RON may pick paths that circumvent policies
- Then again, RON has its own overheads
  - Packets go in and out at intermediate nodes
    - Performance degradation, load on hosts, and financial cost
  - Probing overhead to monitor the virtual links
    - Limits RON to deployments with a small number of nodes

# Peer-to-Peer Networks: Napster



- **Napster history: the rise**

- January 1999: Napster version 1.0
- May 1999: company founded
- September 1999: first lawsuits
- 2000: 80 million users



**Shawn Fanning,  
Northeastern freshman**

- **Napster history: the fall**

- Mid 2001: out of business due to lawsuits
- Mid 2001: dozens of P2P alternatives that were harder to touch, though these have gradually been constrained
- 2003: growth of pay services like iTunes

- **Napster history: the resurrection**

- 2003: Napster reconstituted as a pay service
- 2006: still lots of file sharing going on

# Napster Technology: Directory Service



- User installing the software
  - Download the client program
  - Register name, password, local directory, etc.
- Client contacts Napster (via TCP)
  - Provides a list of music files it will share
  - ... and Napster's central server updates the directory
- Client searches on a title or performer
  - Napster identifies online clients with the file
  - ... and provides IP addresses
- Client requests the file from the chosen supplier
  - Supplier transmits the file to the client
  - Both client and supplier report status to Napster



# Napster Technology: Properties



- Server's directory continually updated
  - Always know what music is currently available
  - Point of vulnerability for legal action
- Peer-to-peer file transfer
  - No load on the server
  - Plausible deniability for legal action (but not enough)
- Proprietary protocol
  - Login, search, upload, download, and status operations
  - No security: cleartext passwords and other vulnerability
- Bandwidth issues
  - Suppliers ranked by apparent bandwidth & response time

# Napster: Limitations of Central Directory



- Single point of failure
- Performance bottleneck
- Copyright infringement

File transfer is decentralized, but locating content is highly centralized

- So, later P2P systems were more distributed

# Peer-to-Peer Networks: Gnutella



- Gnutella history
  - 2000: J. Frankel & T. Pepper released Gnutella
  - Soon after: many other clients (e.g., Morpheus, Limewire, Bearshare)
  - 2001: protocol enhancements, e.g., “ultrapeers”
- Query flooding
  - Join: contact a few nodes to become neighbors
  - Publish: no need!
  - Search: ask neighbors, who ask their neighbors
  - Fetch: get file directly from another node



# Gnutella: Query Flooding



- Fully distributed
  - No central server
- Public domain protocol
- Many Gnutella clients implementing protocol

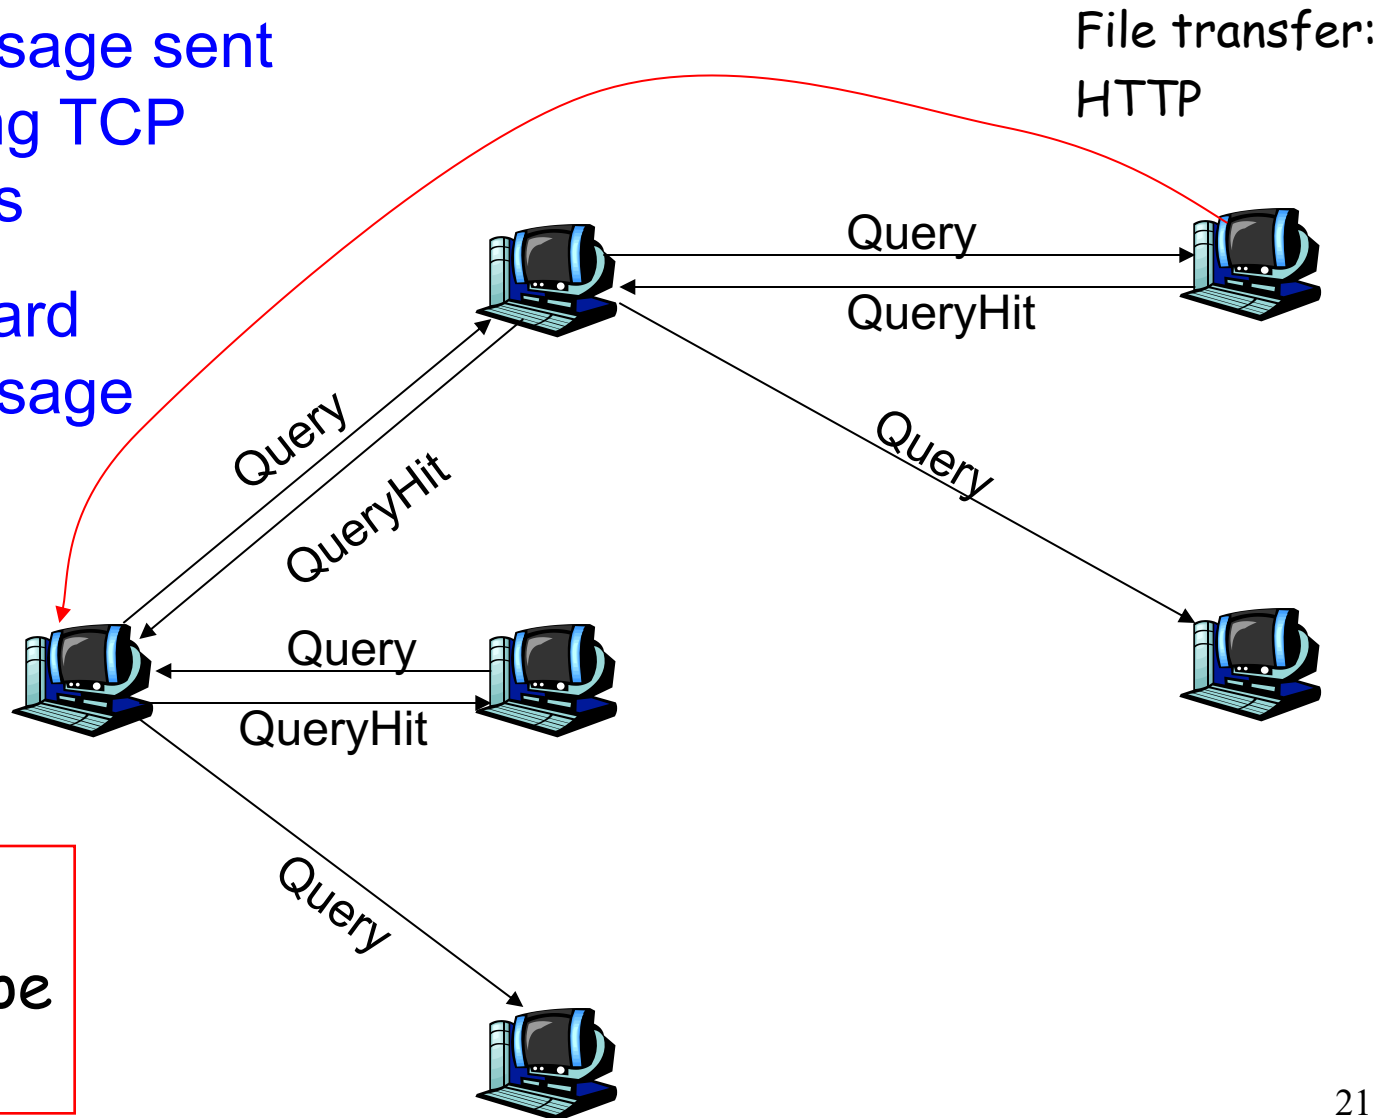
## Overlay network: graph

- Edge between peer X and Y if there's a TCP connection
- All active peers and edges is overlay net
- Given peer will typically be connected with  $< 10$  overlay neighbors

# Gnutella: Protocol



- Query message sent over existing TCP connections
- Peers forward Query message
- QueryHit sent over reverse path



Scalability:  
limited scope  
flooding

# Gnutella: Peer Joining



- Joining peer X must find some other peer in Gnutella network: use list of candidate peers
- X sequentially attempts to make TCP with peers on list until connection setup with Y
- X sends Ping message to Y; Y forwards Ping message.
- All peers receiving Ping message respond with Pong message
- X receives many Pong messages. It can then setup additional TCP connections

# Gnutella: Pros and Cons



- Advantages

- Fully decentralized
- Search cost distributed
- Processing per node permits powerful search semantics

- Disadvantages

- Search scope may be quite large
- Search time may be quite long
- High overhead and nodes come and go often

# Peer-to-Peer Networks: KaAzA



- KaZaA history

- 2001: created by Dutch company (Kazaa BV)
- Single network called FastTrack used by other clients as well
- Eventually the protocol changed so other clients could no longer talk to it



- Smart query flooding


- Join: on start, the client contacts a super-node (and may later become one)
- Publish: client sends list of files to its super-node
- Search: send query to super-node, and the super-nodes flood queries among themselves
- Fetch: get file directly from peer(s); can fetch from multiple peers at once



# KaZaA: Exploiting Heterogeneity



- Each peer is either a group leader or assigned to a group leader
  - TCP connection between peer and its group leader
  - TCP connections between some pairs of group leaders
- Group leader tracks the content in all its children



● ordinary peer

● group-leader peer

— neighboring relationships  
in overlay network

# KaZaA: Motivation for Super-Nodes



- Query consolidation
  - Many connected nodes may have only a few files
  - Propagating query to a sub-node may take more time than for the super-node to answer itself
- Stability
  - Super-node selection favors nodes with high up-time
  - How long you've been on is a good predictor of how long you'll be around in the future

# Peer-to-Peer Networks: BitTorrent



- BitTorrent history and motivation
  - 2002: B. Cohen debuted BitTorrent
  - Key motivation: popular content
    - Popularity exhibits temporal locality (Flash Crowds)
    - E.g., Slashdot effect, CNN Web site on 9/11, release of a new movie or game
  - Focused on efficient *fetching*, not searching
    - Distribute same file to many peers
    - Single publisher, many downloaders
  - Preventing free-loading



# BitTorrent: Simultaneous Downloading



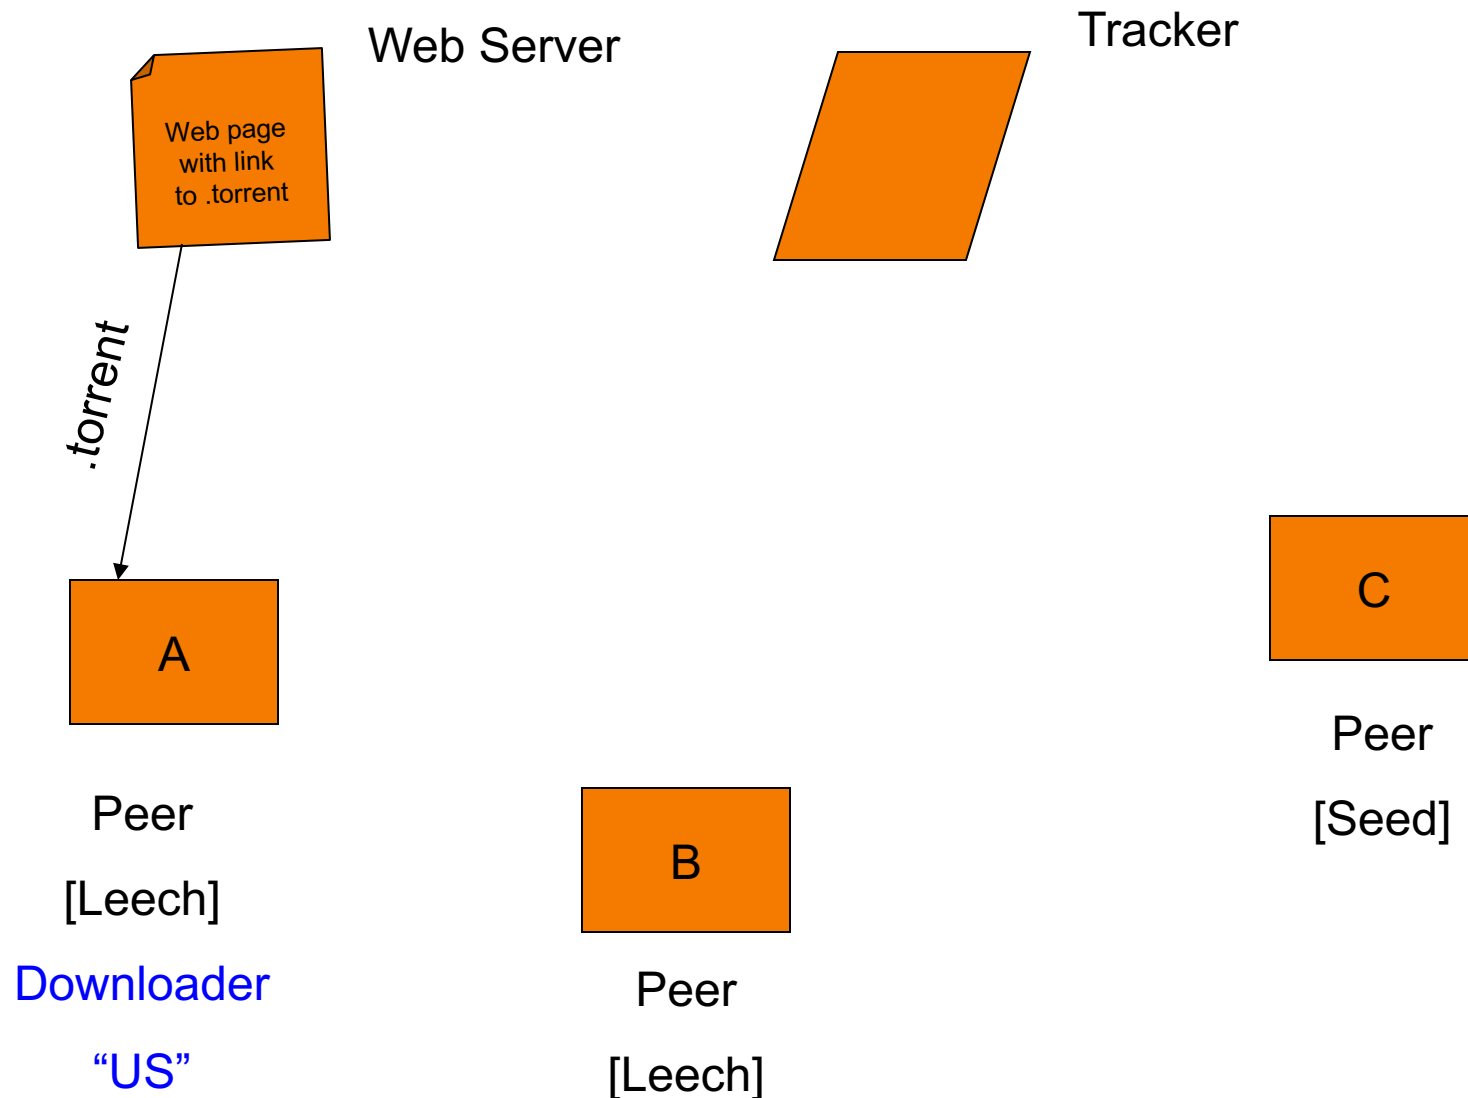
- Divide large file into many pieces
  - Replicate different pieces on different peers
  - A peer with a complete piece can trade with other peers
  - Peer can (hopefully) assemble the entire file
- Allows simultaneous downloading
  - Retrieving different parts of the file from different peers at the same time



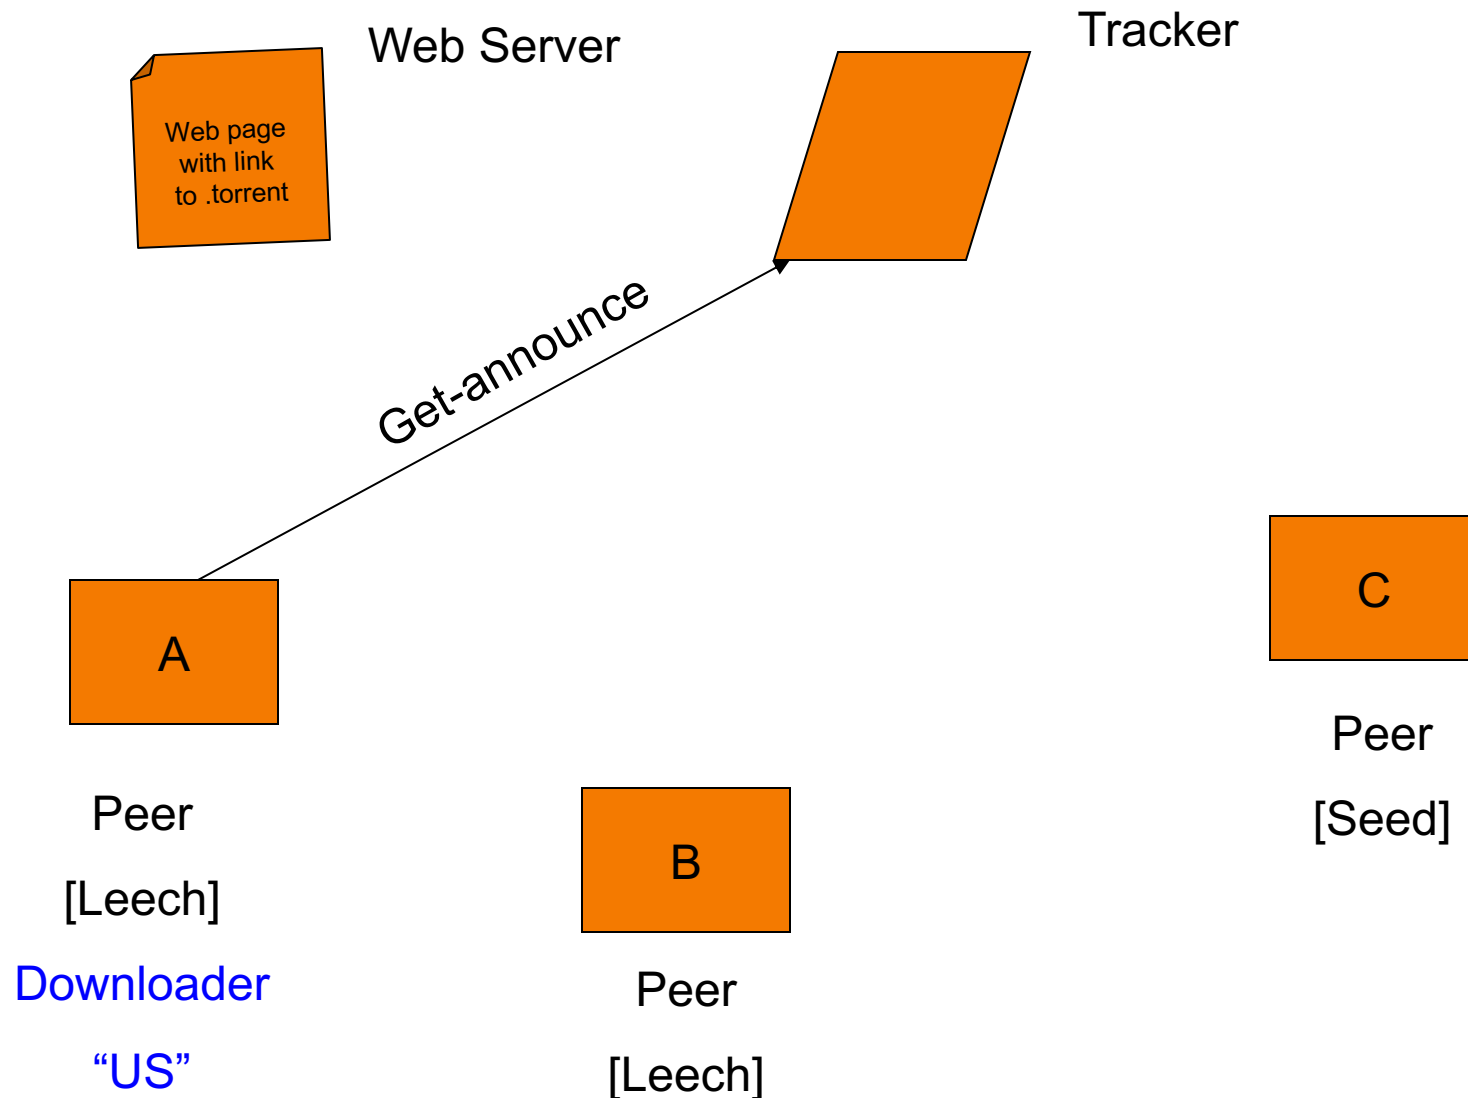
# BitTorrent Components

- Seed
  - Peer with entire file
  - Fragmented in pieces
- Leacher
  - Peer with an incomplete copy of the file
- Torrent file
  - Passive component
  - Stores summaries of the pieces to allow peers to verify their integrity
- Tracker
  - Allows peers to find each other
  - Returns a list of random peers

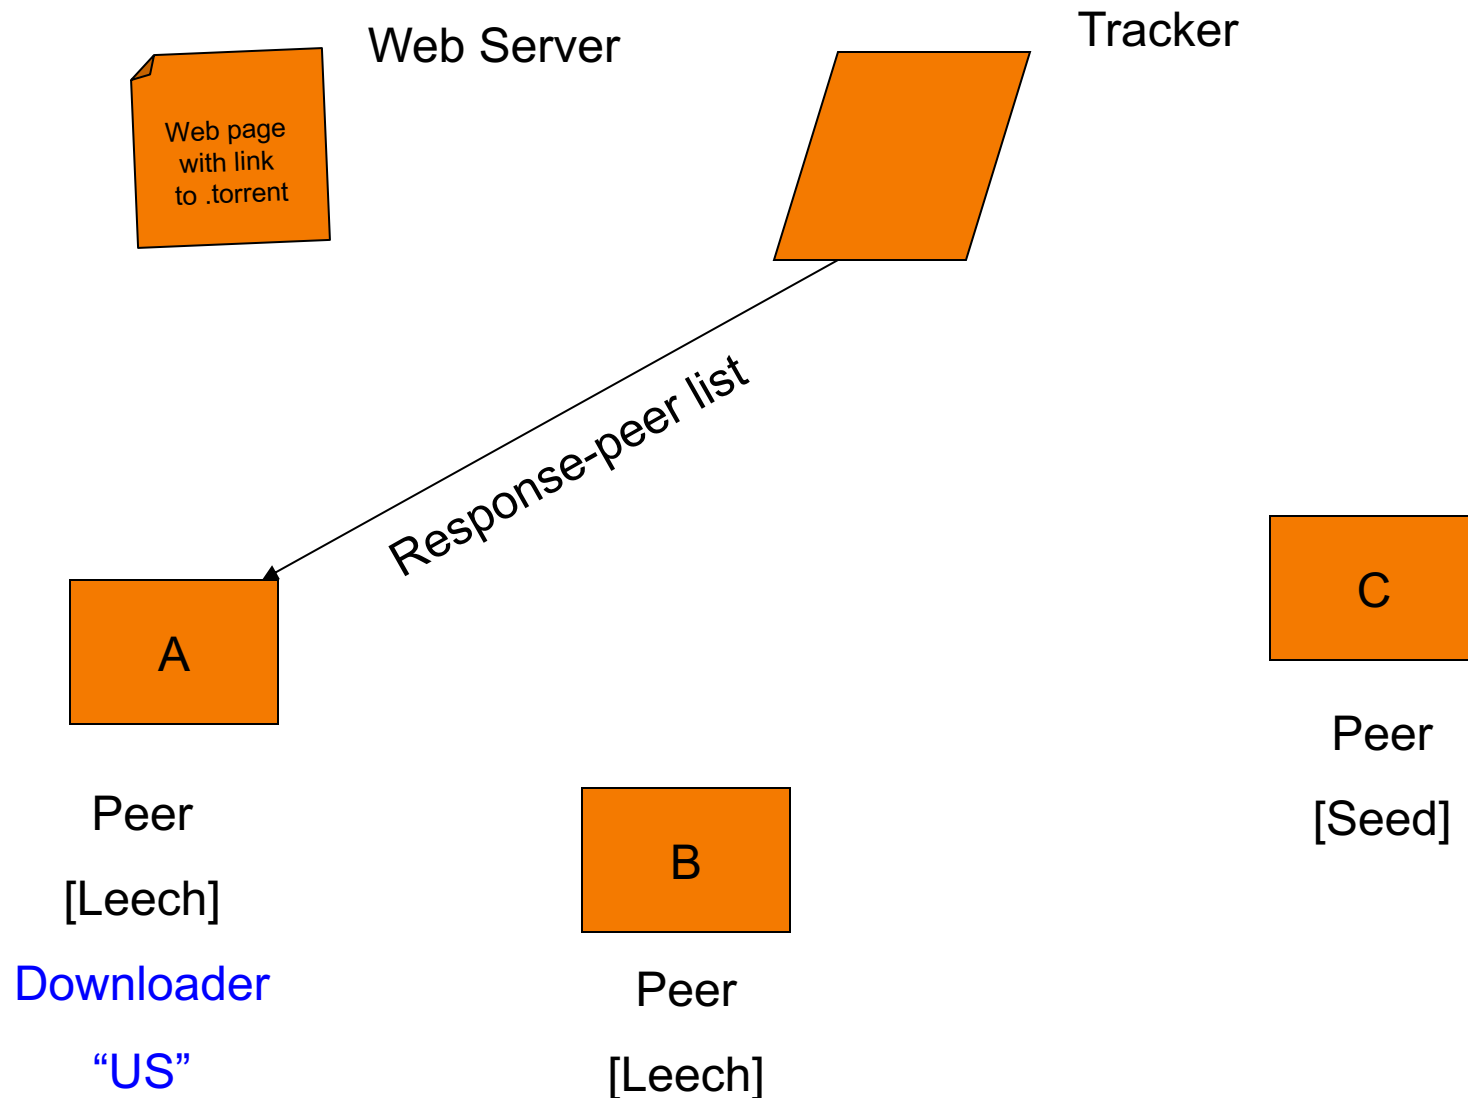
# BitTorrent: Overall Architecture



# BitTorrent: Overall Architecture

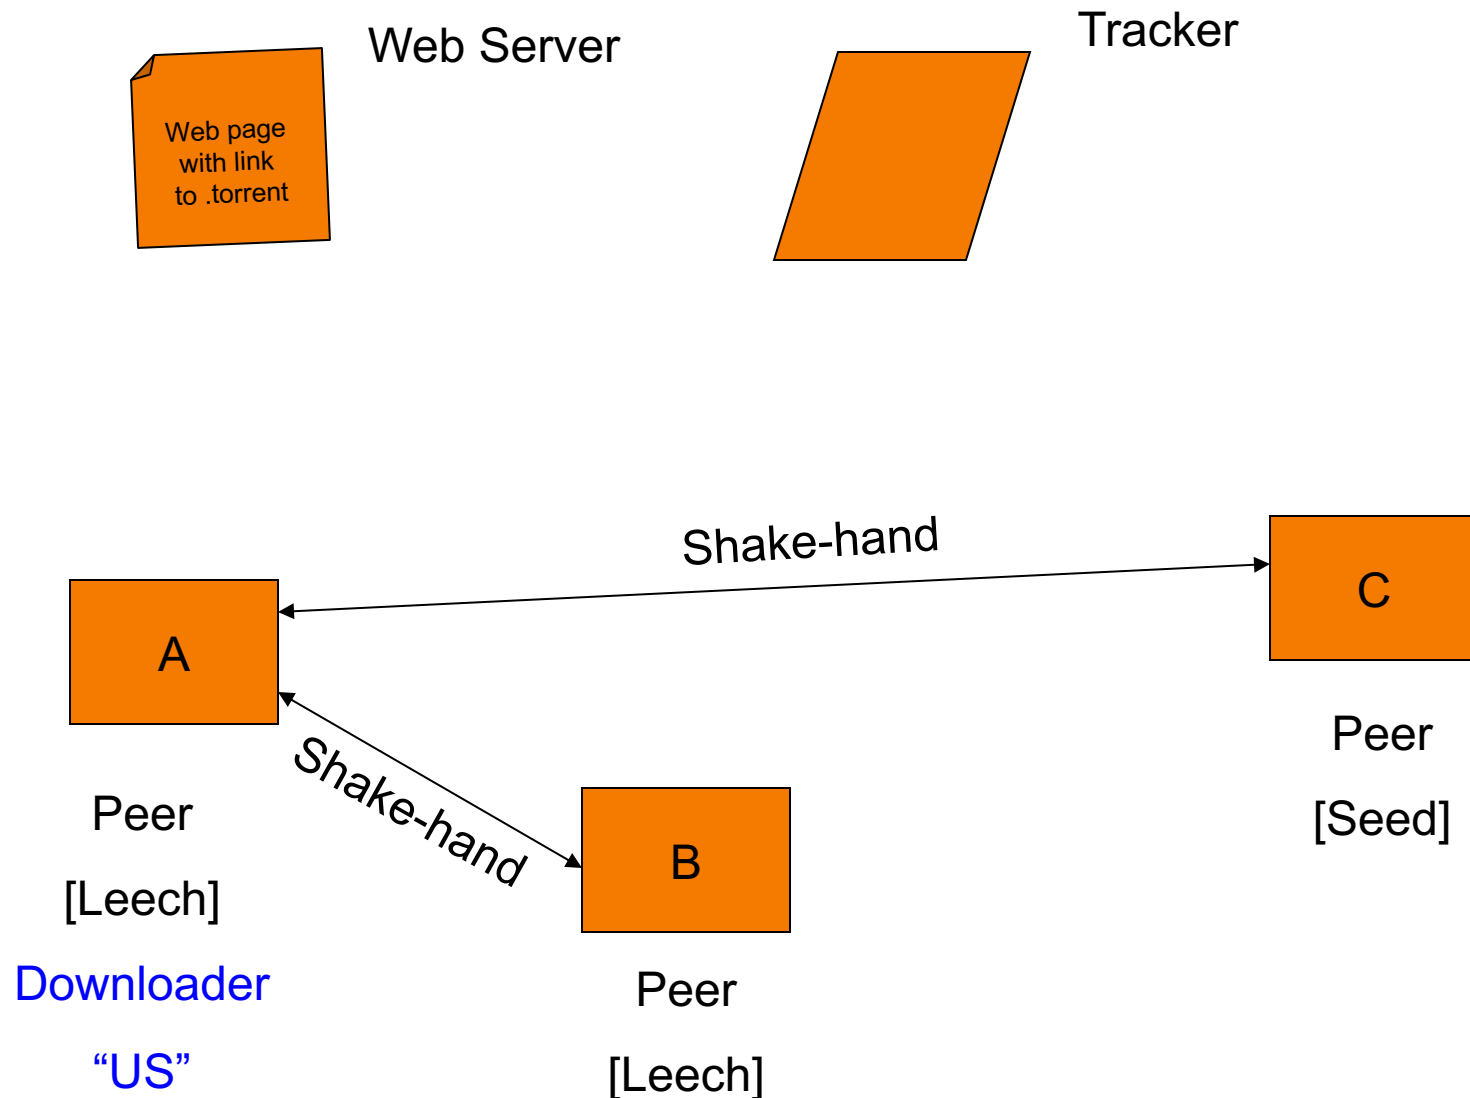


# BitTorrent: Overall Architecture

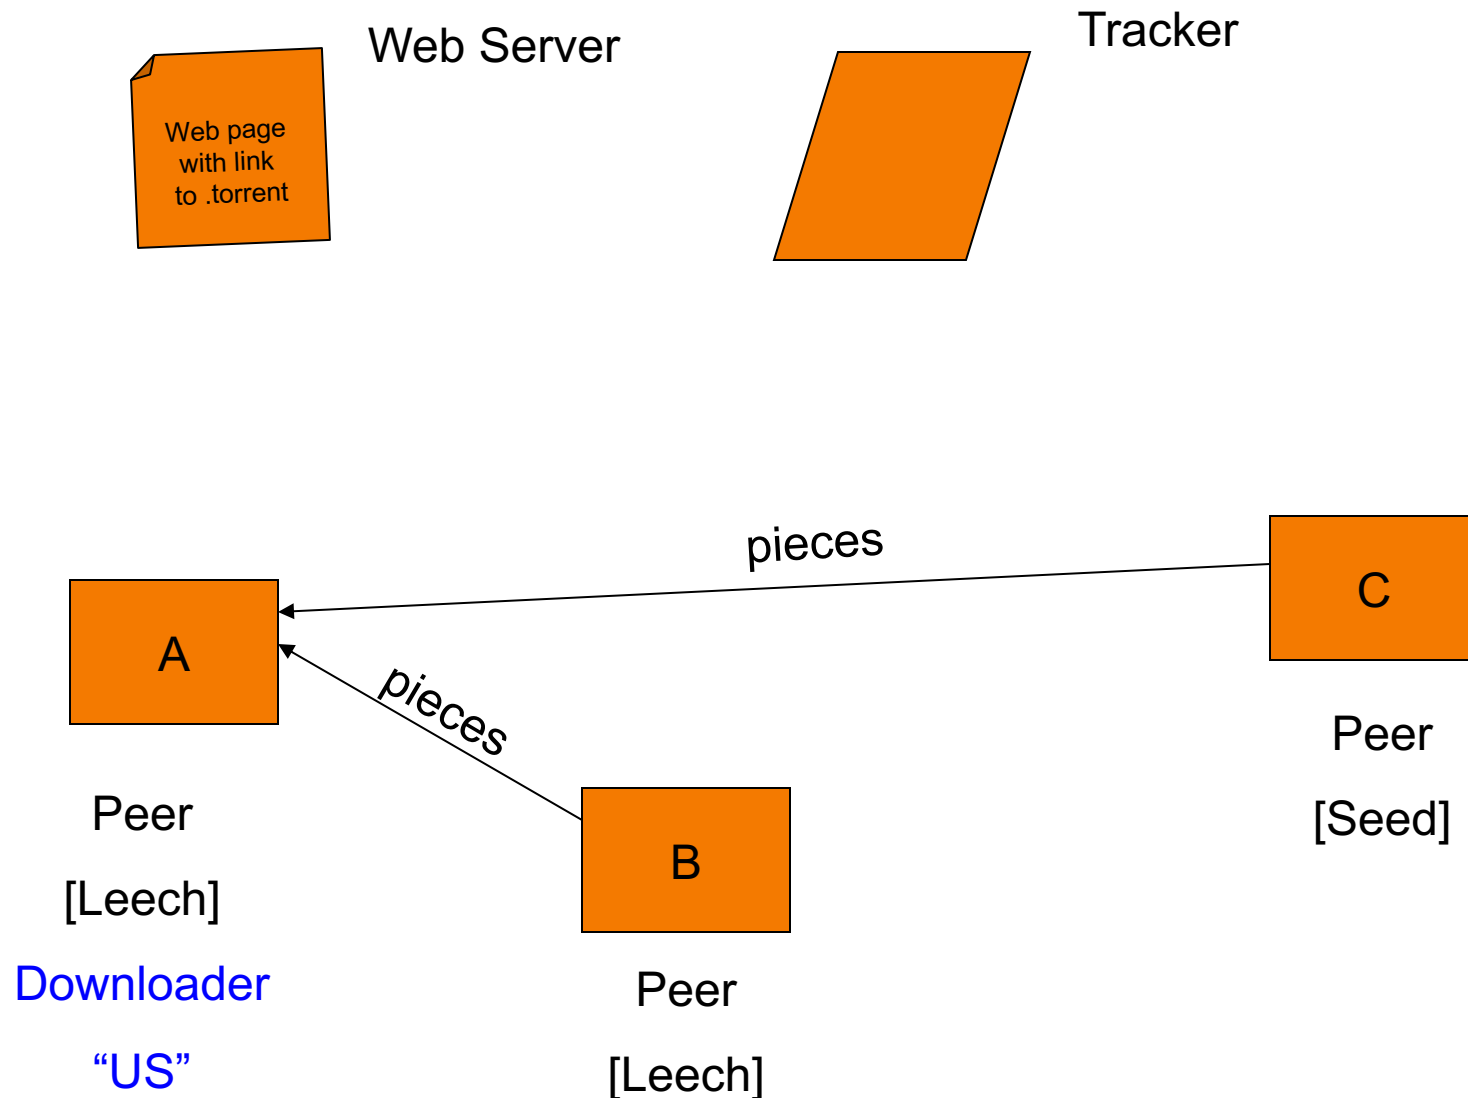




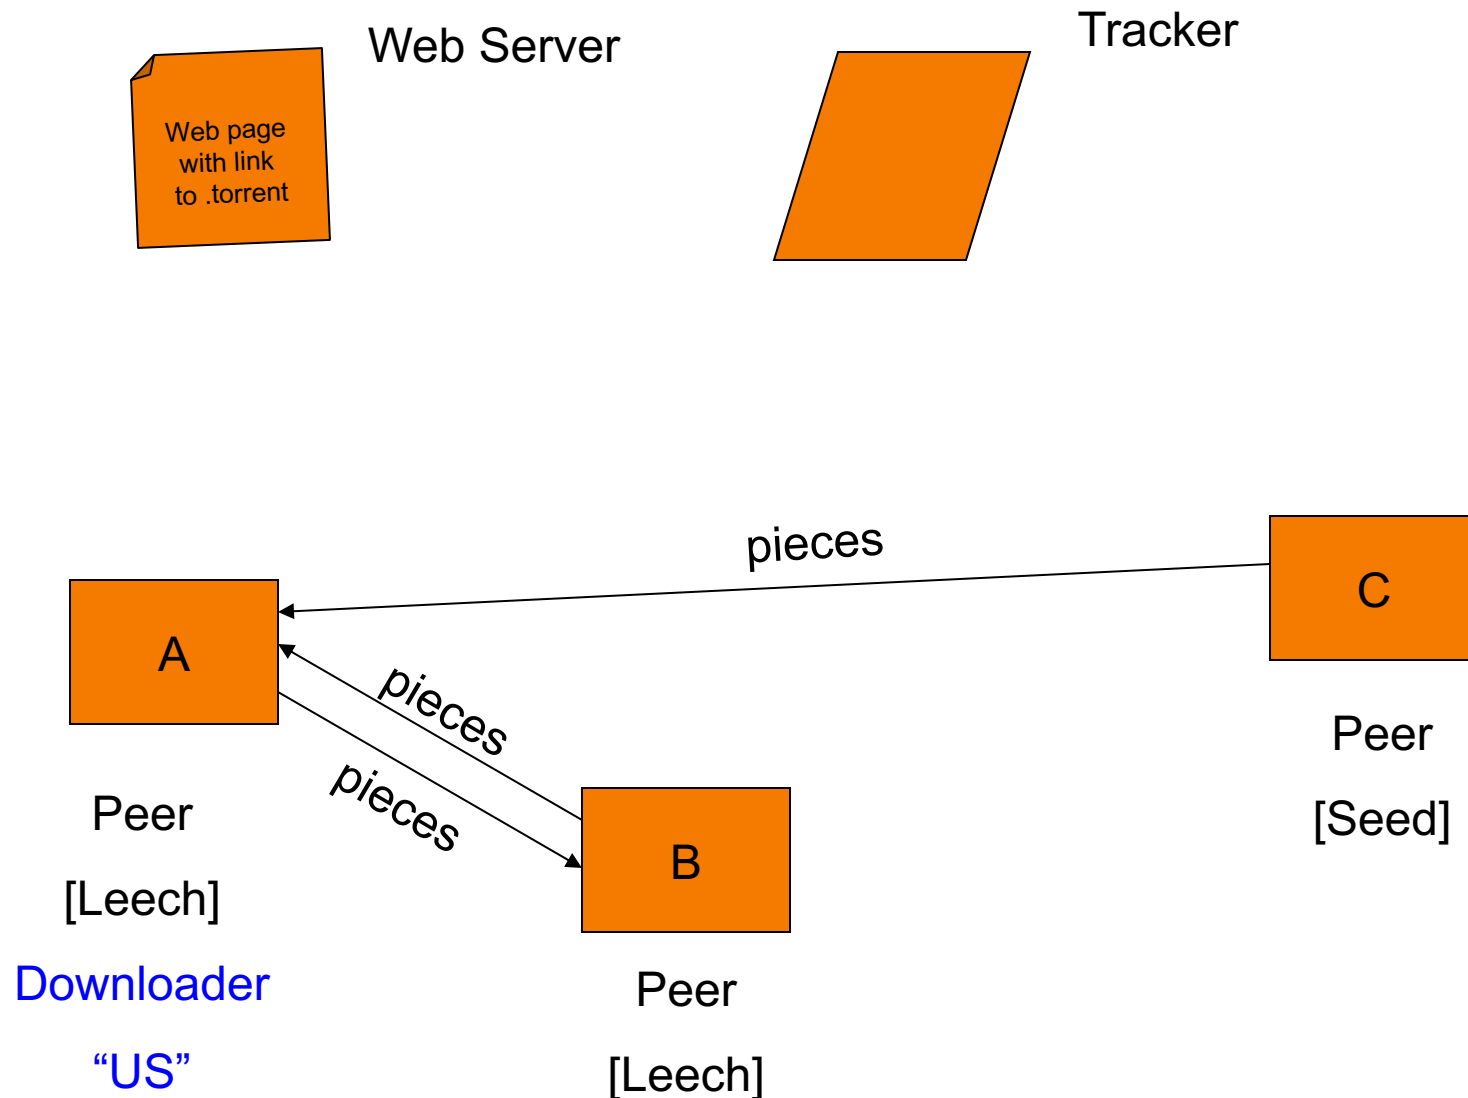
# BitTorrent: Overall Architecture



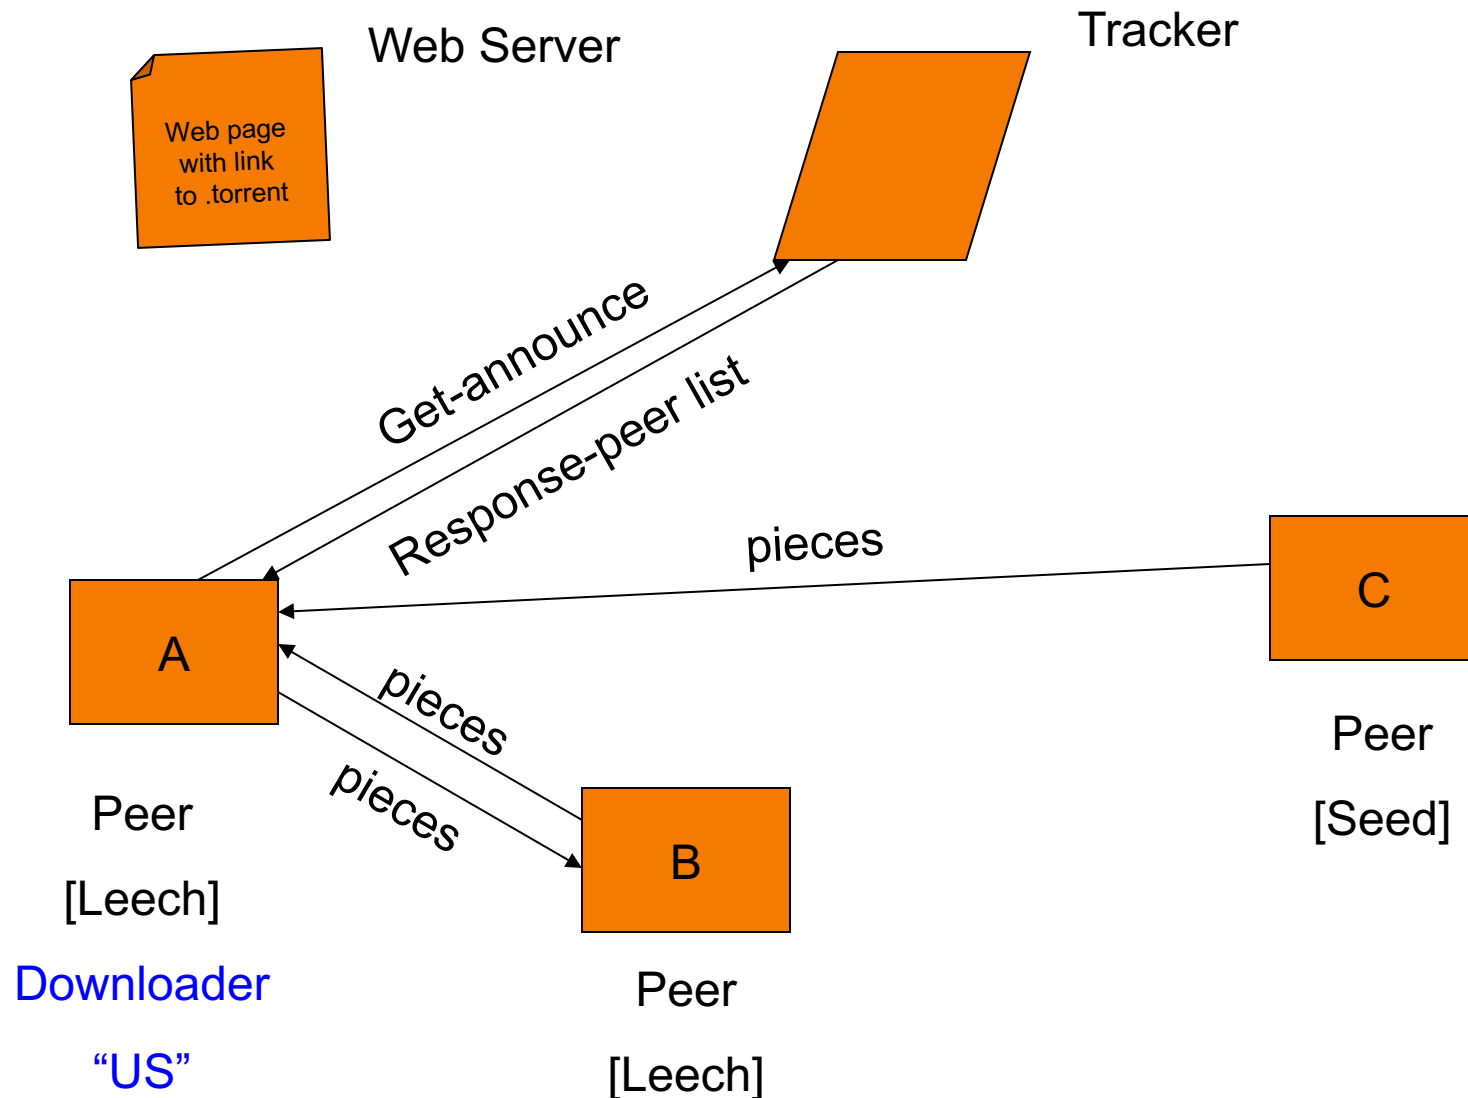
# BitTorrent: Overall Architecture



# BitTorrent: Overall Architecture



# BitTorrent: Overall Architecture



# Free-Riding Problem in P2P Networks



- Vast majority of users are free-riders
  - Most share no files and answer no queries
  - Others limit # of connections or upload speed
- A few “peers” essentially act as servers
  - A few individuals contributing to the public good
  - Making them hubs that basically act as a server
- BitTorrent prevent free riding
  - Allow the fastest peers to download from you
  - Occasionally let some free loaders download

# Conclusions



- **Overlay networks**
  - Tunnels between host computers
  - Hosts implement new protocols and services
  - Effective way to build networks on top of the Internet
- **Peer-to-peer networks**
  - Nodes are end hosts
  - Primarily for file sharing, and recently telephony
  - Centralized directory (Napster), query flooding (Gnutella), super-nodes (KaZaA), and distributed downloading and anti-free-loading (BitTorrent)
- **Great example of how change can happen so quickly in application-level protocols**