

Transport Control Protocol

Outline

- TCP objectives revisited

- TCP basics


- New algorithms for RTO calculation

TCP Overview

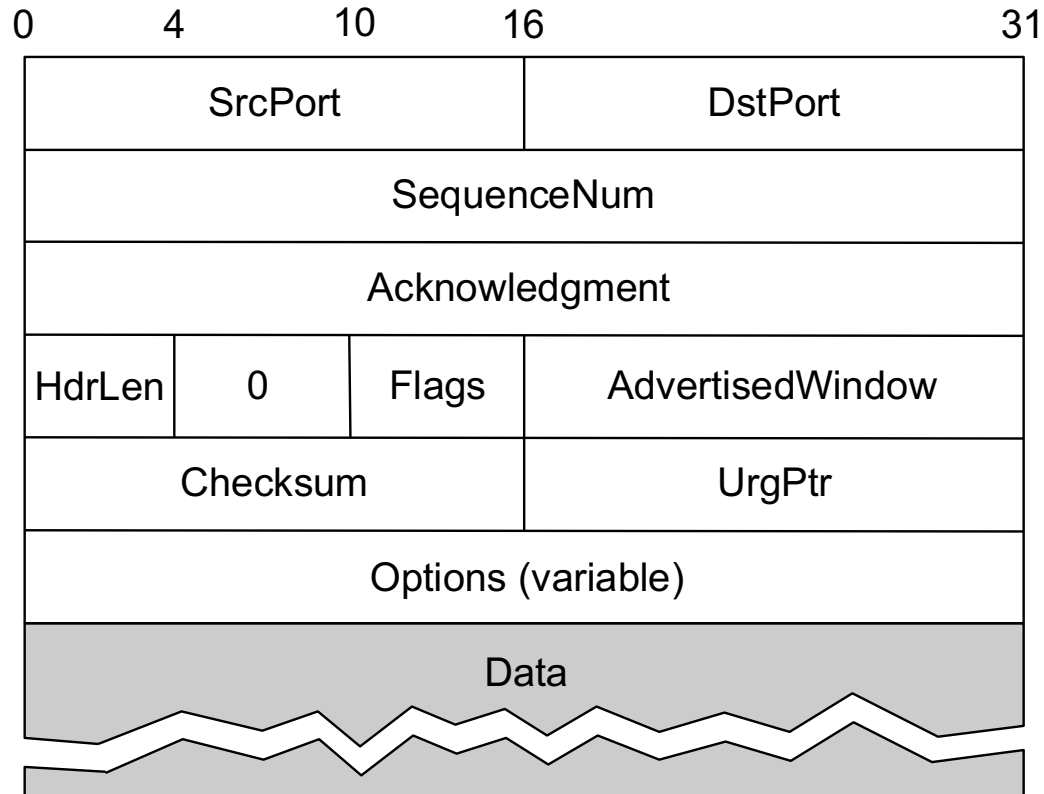
- TCP is the most widely used Internet protocol
 - Web, Peer-to-peer, FTP, telnet, ...
- A two way, reliable, byte stream oriented end-to-end protocol
 - Includes flow and congestion control
- Closely tied to the Internet Protocol (IP)
- A focus of intense study for many years
 - Our goal is to understand the RENO version of TCP
 - RENO is most widely used TCP today
 - RFC 2001 (now expired)
 - RENO mainly specifies mechanisms for dealing with congestion

TCP Features

- Connection-oriented
- Byte-stream
 - app writes bytes
 - TCP sends *segments*
 - app reads bytes
- Reliable data transfer
- Full duplex
- Flow control: keep sender from overrunning receiver
- Congestion control: keep sender from overrunning network

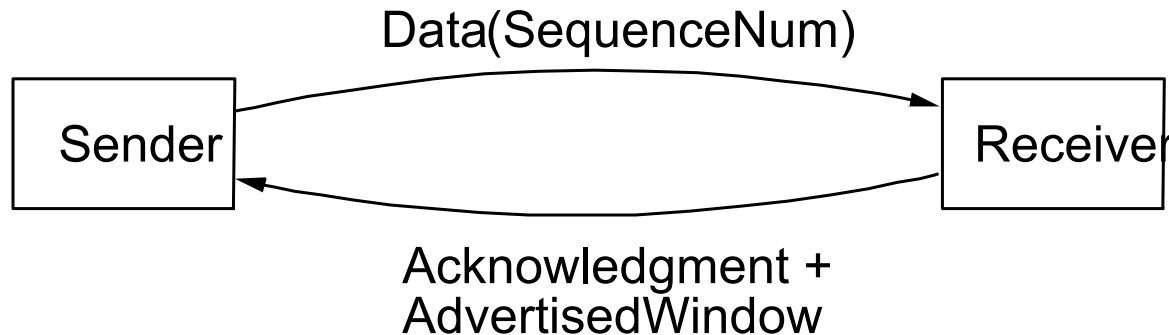


Segment Format



Segment Format (cont)

- Each connection identified with 4-tuple:
 - **(SrcPort, SrcIPAddr, DsrPort, DstIPAddr)**
- Sliding window + flow control
 - **acknowledgment, SequenceNum, AdvertisedWinow**



- Flags
 - **SYN, FIN, RESET, PUSH, URG, ACK**
- Checksum is the same as UDP
 - pseudo header + TCP header + data

Sequence Numbers

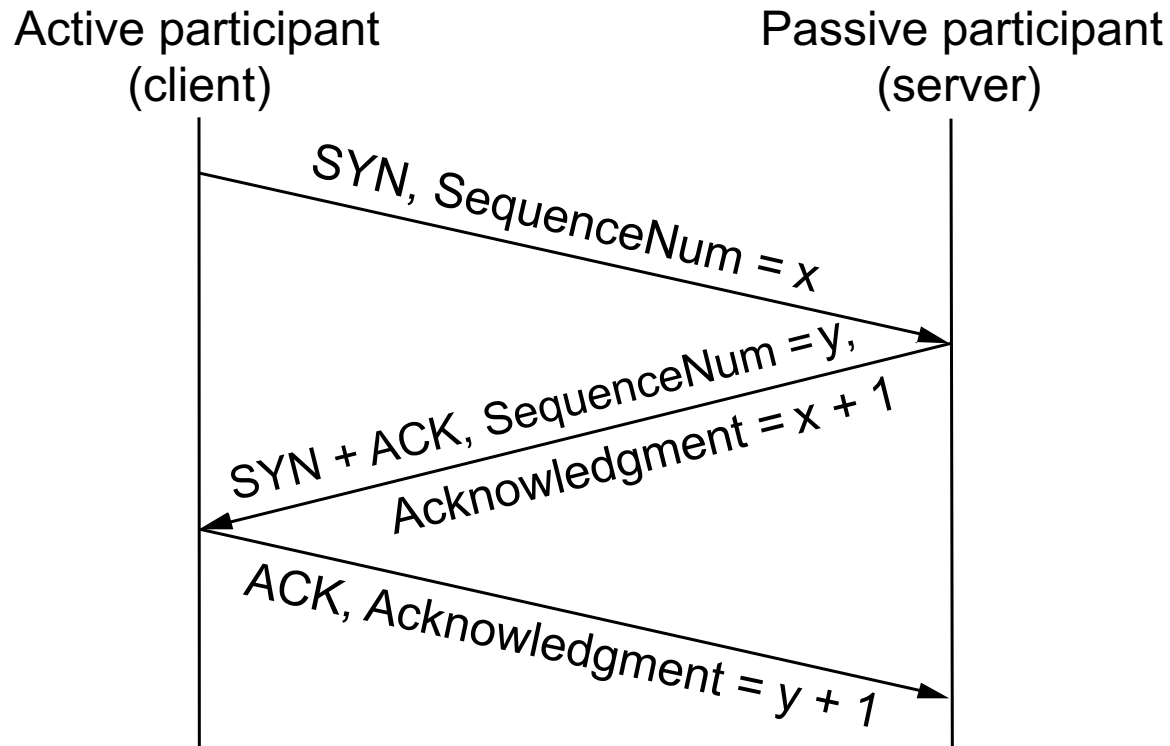
- 32 bit sequence numbers
 - Wrap around supported
- TCP breaks byte stream from application into packets (limited by Max. Segment Size)
- Each byte in the data stream is considered
- Each packet has a sequence number
 - Initial number selected at connection time
 - Subsequent numbers indicate first data byte number in packet
- ACK' s indicate *next byte expected*

Sequence Number Wrap Around

Bandwidth	Time Until Wrap Around
T1 (1.5 Mbps)	6.4 hours
Ethernet (10 Mbps)	57 minutes
T3 (45 Mbps)	13 minutes
FDDI (100 Mbps)	6 minutes
STS-3 (155 Mbps)	4 minutes
STS-12 (622 Mbps)	55 seconds
STS-24 (1.2 Gbps)	28 seconds

- Protect against this by adding a 32-bit timestamp to TCP header

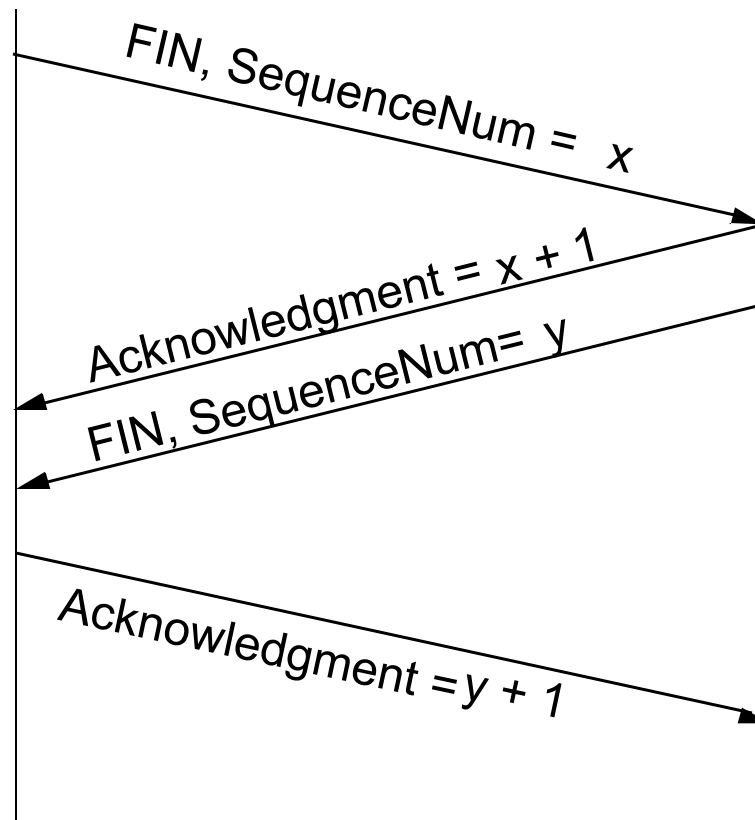
Connection Establishment



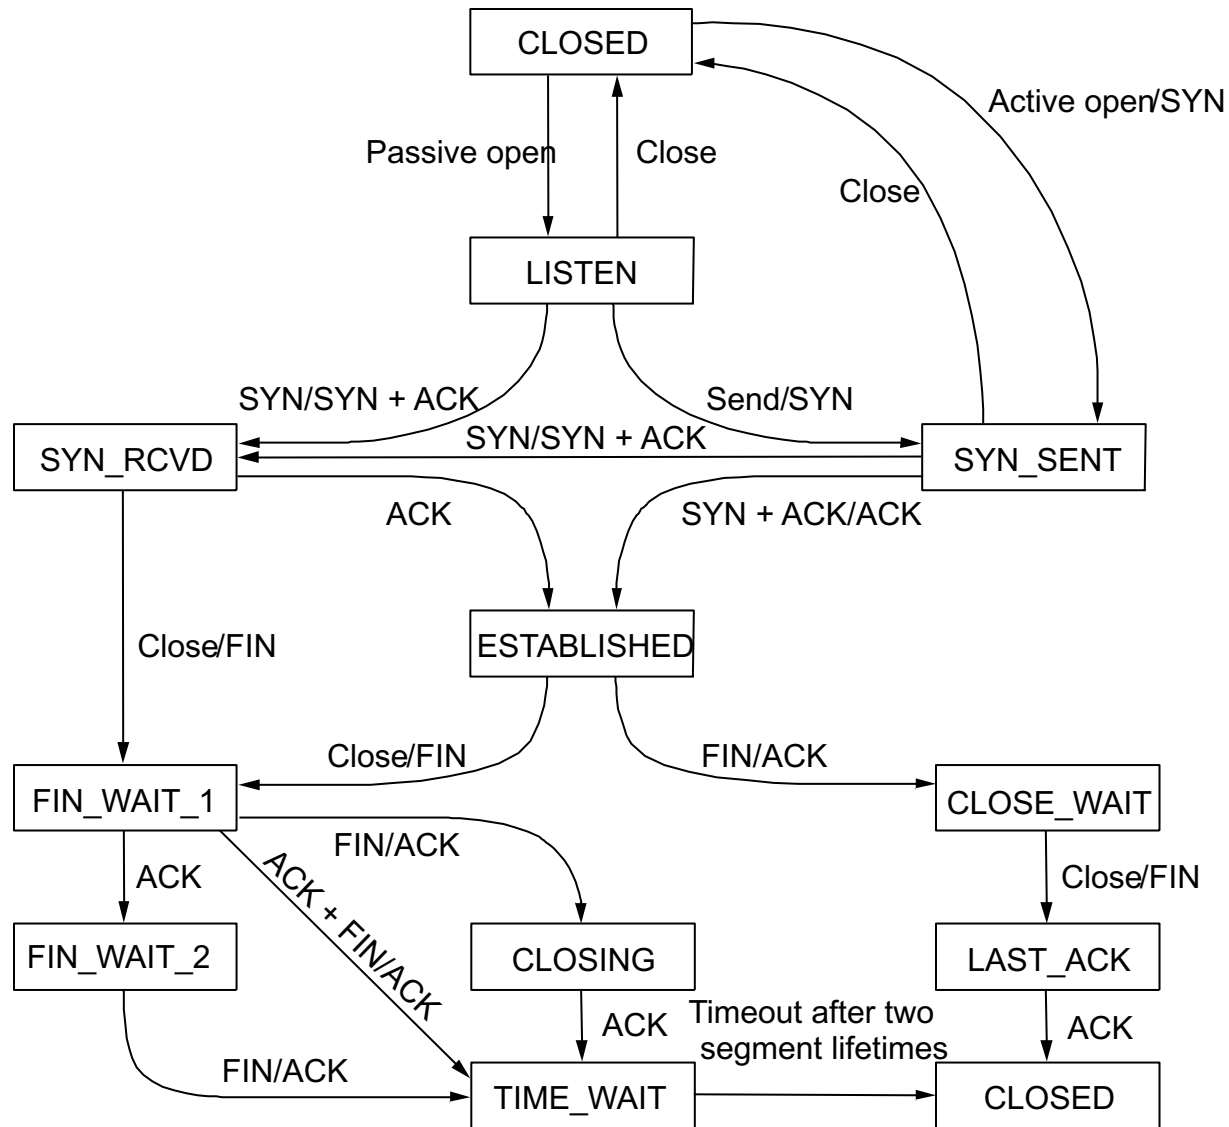
Connection Termination

Active participant
(server)

Passive participant
(client)



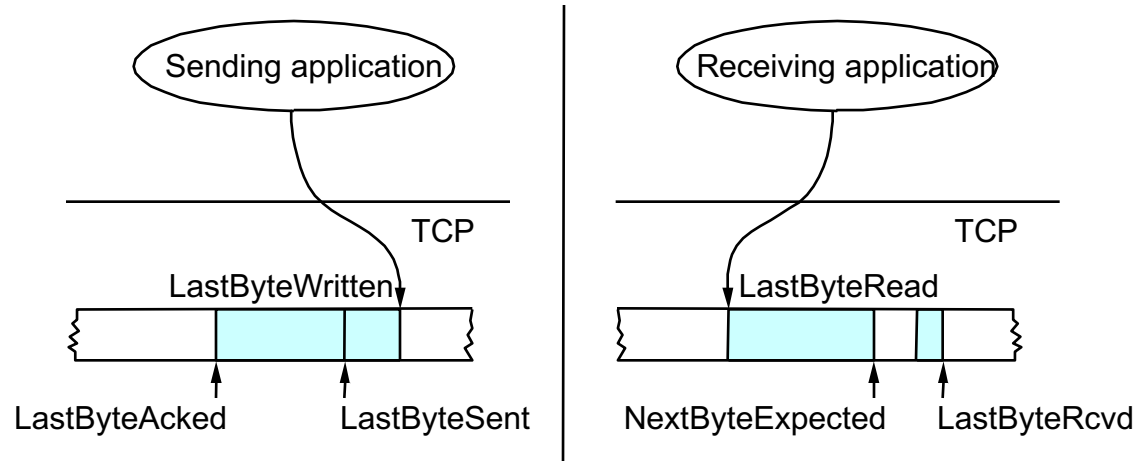
State Transition Diagram



Reliability in TCP

- Checksum used to detect bit level errors
- Sequence numbers used to detect sequencing errors
 - Duplicates are ignored
 - Reordered packets are reordered (or dropped)
 - Lost packets are retransmitted
- Timeouts used to detect lost packets
 - Requires RTO calculation
 - Requires sender to maintain data until it is ACKed

Sliding Window Revisited



- Sending side

- **LastByteAcked** \leq **LastByteSent**
- **LastByteSent** \leq **LastByteWritten**
- buffer bytes between **LastByteAcked** and **LastByteWritten**

- Receiving side

- **LastByteRead** $<$ **NextByteExpected**
- **NextByteExpected** \leq **LastByteRcvd** + 1
- buffer bytes between **NextByteRead** and **LastByteRcvd**

Flow Control in TCP

- Send buffer size: **MaxSendBuffer**
- Receive buffer size: **MaxRcvBuffer**
- Receiving side
 - $\text{LastByteRcvd} - \text{LastByteRead} \leq \text{MaxRcvBuffer}$
 - $\text{AdvertisedWindow} = \text{MaxRcvBuffer} - (\text{LastByteRcvd} - \text{LastByteRead})$
- Sending side
 - $\text{LastByteSent} - \text{LastByteAcked} \leq \text{AdvertisedWindow}$
 - $\text{EffectiveWindow} = \text{AdvertisedWindow} - (\text{LastByteSent} - \text{LastByteAcked})$
 - $\text{LastByteWritten} - \text{LastByteAcked} \leq \text{MaxSendBuffer}$
 - block sender if $(\text{LastByteWritten} - \text{LastByteAcked}) + y > \text{MaxSenderBuffer}$
- Always send ACK in response to arriving data segment
- Persist sending one byte seg. when **AdvertisedWindow = 0**

Keeping the Pipe Full


- 16-bit **AdvertisedWindow** controls amount of pipelining
- Assume RTT of 100ms
- Add scaling factor extension to header to enable larger windows

Bandwidth	Delay x Bandwidth Product
T1 (1.5 Mbps)	18KB
Ethernet (10 Mbps)	122KB
T3 (45 Mbps)	549KB
FDDI (100 Mbps)	1.2MB
OC-3 (155 Mbps)	1.8MB
OC-12 (622 Mbps)	7.4MB
OC-24 (1.2 Gbps)	14.8MB


Making TCP More Efficient

- Delayed acknowledgements
 - Delay for about 200ms
 - Try to piggyback ACKs with data
- Acknowledge every other packet
 - Many instances in transmission sequence which require an ACK
- Don't forget Nagle's algorithm
 - Can be switched off

Karn/Partridge Algorithm for RTO



- Two degenerate cases with timeouts and RTT measurements
 - Solution: Do not sample RTT when retransmitting
- After each retransmission, set next RTO to be double the value of the last
 - Exponential backoff is well known control theory method
 - Loss is most likely caused by congestion so be careful



Jacobson/ Karels Algorithm

- In late ' 80s, Internet was suffering from *congestion collapse*
- New Calculations for average RTT – Jacobson ' 88
- Variance is not considered when setting timeout value
 - If variance is small, we could set $RTO = EstRTT$
 - If variance is large, we may need to set $RTO > 2 \times EstRTT$
- New algorithm calculates both variance and mean for RTT
- **$Diff = sampleRTT - EstRTT$**
- **$EstRTT = EstRTT + (d \times Diff)$**
- **$Dev = Dev + d (|Diff| - Dev)$**
 - Initially settings for **$EstRTT$** and **Dev** will be given to you
 - where d is a factor between 0 and 1
 - typical value is 0.125

Jacobson/ Karels contd.

- **TimeOut** = $\mu \times \text{EstRTT} + \phi \times \text{Dev}$
 - where $\mu = 1$ and $\phi = 4$
- When variance is small, TimeOut is close to EstRTT
- When variance is large Dev dominates the calculation
- Another benefit of this mechanism is that it is very efficient to implement in code (does not require floating point)
- Notes
 - algorithm only as good as granularity of clock (500ms on Unix)
 - accurate timeout mechanism important to congestion control (later)
- These issues have been studied and dealt with in new RFC's for RTO calculation.
- TCP RENO uses Jacobson/Karels